

Day 6 – Phase 6: Log Rotation, Scheduling, Archiving

Boss's Request: Prepare the system for production use.

Tasks:

- Configure log rotation for `temperature.log` (rotate at 1 MB, compress).
 - **Log rotation** is the process of automatically **managing log files** so they don't grow too large and fill up your disk.
 - I will create configuration file to temperature log rotation
`/etc/logrotate.d/temperature`

```
mohamed@iot ~/iot_logger> cat /etc/logrotate.d/temperature
/home/mohamed/iot_logger/logs/temperature.log {
    size 1M
    rotate 5
    compress
    missingok
    notifempty
    copytruncate
    su mohamed mohamed
}
```

- **size 1M** → rotate when file > 1M
- **rotate 2** → keep 5 old log files
- **compress** → compress rotated logs (`.gz`).
- **missingok** → don't throw error if log is missing.
- **notifempty** → don't rotate if empty.
- **copytruncate** → truncate original log after copying (so process keeps writing) (Don't understand well)
- **su mohamed mohamed** → which user & group to use

- Test by forcing a rotation.
 - At first `logrotate` refuses to rotate files that have **hard links**, because renaming or truncating one would break consistency (I knew that using `sudo logrotate -d /etc/logrotate.d/temperature`)
 - So we will first remove hard link then test

```
rm logs/hard_temperature.log
sudo logrotate -f /etc/logrotate.d/temperature
ls -lah logs/
```

```
mohamed@iot ~/iot_logger> rm logs/hard_temperature.log
mohamed@iot ~/iot_logger> sudo logrotate -f /etc/logrotate.d/temperature
mohamed@iot ~/iot_logger> ls -lah logs/
total 164K
drwxrw---- 2 mohamed mohamed 4.0K Sep  3 18:24 .
drwxrwxr-x 5 mohamed mohamed 4.0K Sep  3 17:03 ..
-rw-rw-r-- 1 mohamed mohamed 21K Sep  3 01:48 filtered.log
lrwxrwxrwx 1 mohamed mohamed 15 Aug 31 21:53 symbolic_temperature.log -> temperature.log
-rwxrw---- 1 mohamed mohamed 2.7K Sep  3 18:24 temperature.log
-rwxrw---- 1 mohamed mohamed 127K Sep  3 18:24 temperature.log.1.gz
-rw-rw-r-- 1 mohamed mohamed  0 Sep  1 16:27 test.txt
mohamed@iot ~/iot_logger> sudo logrotate -f /etc/logrotate.d/temperature
mohamed@iot ~/iot_logger> ls -lah logs/
total 168K
drwxrw---- 2 mohamed mohamed 4.0K Sep  3 18:24 .
drwxrwxr-x 5 mohamed mohamed 4.0K Sep  3 17:03 ..
-rw-rw-r-- 1 mohamed mohamed 21K Sep  3 01:48 filtered.log
lrwxrwxrwx 1 mohamed mohamed 15 Aug 31 21:53 symbolic_temperature.log -> temperature.log
-rwxrw---- 1 mohamed mohamed 559 Sep  3 18:24 temperature.log
-rwxrw---- 1 mohamed mohamed 534 Sep  3 18:24 temperature.log.1.gz
-rwxrw---- 1 mohamed mohamed 127K Sep  3 18:24 temperature.log.2.gz
-rw-rw-r-- 1 mohamed mohamed  0 Sep  1 16:27 test.txt
```

- Schedule the Python script to run every 5 minutes with `cron`.
 - Crontab is a schedule file where you tell Linux: "Run this command at this time." automatically
 - First I will rewrite `~/iot_logger/scripts/sensor_script.py` without `while` or `sleep`

```
import os, time, random

# Get sensor type from environment variable
sensor = os.getenv("SENSOR_TYPE", "unknown")

value = random.randint(15, 45)
print(f"{time.ctime()} | {sensor}: {value}")
```

- Open crontab for current user → `crontab -e`
- Add your schedule task
 - Every 5 minutes run `sensor_script.py` and append sensor data to `temperature.log`
- Make sure every thing is ok → `crontab -l`

```
mohamed@iot ~> crontab -l | tail -n 1
*/5 * * * * SENSOR_TYPE=temperature /usr/bin/python3 /home/mohamed/iot_logger
/scripts/sensor_script.py >> /home/mohamed/iot_logger/logs/temperature.log
```

- Check `temperature.log` files

```
mohamed@iot ~/iot_logger> tail -n 3 logs/temperature.log
Wed Sep  3 20:35:01 2025 | temperature: 15
Wed Sep  3 20:40:01 2025 | temperature: 38
Wed Sep  3 20:45:01 2025 | temperature: 36
```

- Verify log growth over time.

```
mohamed@iot ~/iot_logger> ls -lah logs/temperature.log
-rwxrw---- 1 mohamed mohamed 190K Sep  3 20:40 logs/temperature.log
```

```
mohamed@iot ~/iot_logger> ls -lah logs/temperature.log
-rwxrw---- 1 mohamed mohamed 191K Sep  3 22:05 logs/temperature.log
```

- Compress old logs into `.tar.gz` in `data/`.

```
tar -czf data/old_logs.tar.gz logs/*.gz
```

```
mohamed@iot ~/iot_logger> ls -lh data/old_logs.tar.gz
-rw-rw-r-- 1 mohamed mohamed 311K Sep  3 20:57 data/old_logs.tar.gz
mohamed@iot ~/iot_logger> tar -tzf data/old_logs.tar.gz
logs/temperature.log.1.gz
logs/temperature.log.2.gz
logs/temperature.log.3.gz
```

- Simulate sending archives to `/home/<username>/server/` using `cp`, `scp`, or `rsync`.
(hint: you can use `scp` and `copy` to destination directory in another path on the same machine just for simulation).
 - Create `server` folder for simulation → `mkdir ~/server`
 - Simulate remote copy → `scp ~/iot_logger/data/*.tar.gz ~/server/`

```
mohamed@iot ~> mkdir ~/server
mohamed@iot ~> ls server/
mohamed@iot ~> scp ~/iot_logger/data/*.tar.gz ~/server/
mohamed@iot ~> ls server/
old_logs.tar.gz
```

- You can also use `cp`, `rsync` instead of `scp`
 - `cp` → Local system only
 - `rsync` → Local and remote
 - `scp` → Local ↔ Remote

Open-Ended Questions:

- How does `cron` scheduling work? Show a crontab entry to run a script every 5 minutes.
 - **Cron** is a background service that runs jobs at scheduled times.
 - **Cron** is a background service that runs jobs at scheduled times.
 - Example:

```
* /5 * * * * /usr/bin/python3 /home/mohamed/iot_logger/scripts/sensor_script.py >> /home/mohamed/iot_logger/logs/temperature.log
```

- Why do we need log rotation? Show an example logrotate config for temperature.log.
 - Logs grow continuously and can fill up disk space.
 - **Log rotation** keeps logs manageable
 - Example:

```
/home/mohamed/iot_logger/logs/temperature.log {  
    size 1M  
    rotate 5  
    compress  
    missingok  
    notifempty  
    copytruncate  
    su mohamed mohamed  
}
```

- Explain the difference between a Virtual Machine and a Container. Must containers use the same OS as the host? Why or why not?
 - **Virtual Machine (VM):**
 - Runs full OS with its **own kernel**.
 - Heavyweight → more CPU & RAM usage.
 - Strong isolation (hardware-level).
 - Slow startup (minutes).
 - **Container:**
 - **Shares** host's **kernel**.
 - Lightweight → less resource usage.
 - Isolation at process-level.
 - Fast startup (seconds).
 - **Do containers need same OS?**
 - Must share the same kernel family with different distributions.
- Reflection: Which actions in this project combined multiple Linux concepts (e.g., redirection + process monitoring)? How does this apply to real IoT systems?
 - I think `python3 script.py >> logs/temperature.log &`
 - `>>` → redirection
 - `&` → process monitoring
 - And `logrotate` & `cron` → file management , compression, scheduling
 - Real **IoT** systems:
 - Need rotation to avoid storage overflow.
 - Scheduling ensures periodic data collection.
 - Archiving + transfer → central monitoring system.